

Unsupervised Machine Learning - Analysis Report

1. Dataset Generation

A synthetic classification dataset was generated using `sklearn.datasets.make_classification` with the following configuration:


Parameter	Value	Rationale
<code>n_samples</code>	200-1000 (random)	Within required range
<code>n_features</code>	2	Enables 2D visualisation
<code>n_informative</code>	2	All features carry signal
<code>n_redundant</code>	0	No artificially correlated features
<code>n_classes</code>	4	Maximum with 2 informative features (see below)
<code>n_clusters_per_class</code>	1	One Gaussian blob per class
<code>flip_y</code>	0.05	5 % label noise for realism
<code>random_state</code>	42	Reproducibility

Why 4 classes and not more?

`make_classification` enforces the constraint: `n_classes × n_clusters_per_class ≤ 2n_informative`.

With 2 informative features that ceiling is **2² = 4**, making 4 the maximum number of separable classes when

`n_clusters_per_class = 1`.



Features were standardised with `StandardScaler` before clustering. KMeans is a distance-based algorithm - without standardisation, a feature with a larger numeric range would dominate the distance calculation and bias the cluster boundaries.

2. KMeans Model Fitting

A baseline KMeans model was fitted with **k = 4** (matching the number of true classes).

```
KMeans(n_clusters=4, n_init=10, random_state=42)
```


Using `n_init=10` (the default) means the algorithm runs 10 independent initialisations and keeps the solution with the lowest inertia (within-cluster sum of squared distances). This makes the result robust to unlucky centroid starts.

3. Decision-Region Visualisation (Meshgrid Plot)

The fitted model was visualised using a **Voronoi / decision-boundary plot**:

1. A dense grid of points (`h = 0.02` step) was generated to cover the entire feature space.
2. Every grid point was assigned a cluster label by `model.predict()`.
3. `contourf` shaded the resulting regions, one colour per cluster.
4. Actual data points were overlaid, coloured by their predicted cluster.
5. Centroids were marked with a star (★).

This makes the cluster boundaries and centroid positions directly interpretable alongside the data distribution.



4. Silhouette Score

The **silhouette score** quantifies how well each point fits its assigned cluster relative to the nearest alternative cluster.

For a single point i , the score is:

$$s(i) = (b(i) - a(i)) / \max(a(i), b(i))$$

where $a(i)$ is the mean intra-cluster distance and $b(i)$ is the mean distance to the nearest neighbouring cluster. The overall score is the average across all points.

Score range Interpretation

close to **+1** Point is well inside its cluster and far from others

close to **0** Point sits on or near a cluster boundary

close to **-1** Point is closer to a neighbouring cluster than its own

```
silhouette_score(X, kmeans.labels_)
```

```
# Baseline result (k=4, seed=42): reported in notebook output
```

5. Sensitivity Analysis

All experiments in this section use **n_init=1**. With a single initialisation, seed-level randomness is not averaged away, which deliberately exposes how much the result depends on the chosen k and the random starting positions.

5.1 Effect of k on the Silhouette Score

Three k -values were tested with a fixed `random_state=42` : **$k \in \{2, 4, 7\}$** .

k Relationship to true structure

Expected silhouette behaviour

Under-clustering - 4
2 true groups forced
into 2 clusters

Low. Points from structurally different groups share a cluster. Intra-cluster distances increase, silhouette falls.

4

k Relationship to true structure

Matches true class count


Over-clustering -
7 some true groups are split

Expected silhouette behaviour

Highest. The algorithm has the correct number of segments to align with the underlying distribution; intra-cluster distances are minimised without unnecessary fragmentation.

Drops again. Artificial boundaries cut through homogeneous regions. Cohesion increases locally, but separation decreases and interpretability is lost.

Key insight: The silhouette score forms a curve with a peak near the natural cluster count. It is one of the standard heuristics - alongside the elbow method on inertia - for choosing k without ground-truth labels. Neither metric is definitive on its own; combining both gives a more confident decision.



5.2 Effect of Random Seed on the Silhouette Score


Three seeds were tested with fixed $k = 4$ and $n_{\text{init}}=1$: seeds $\in \{0, 42, 99\}$.

KMeans with `k-means++` initialisation places the first centroid uniformly at random, then places subsequent centroids with probability proportional to their squared distance from the nearest already-chosen centroid. Even with `k-means++`, the outcome with $n_{\text{init}}=1$ depends on which point is selected as the first centroid - controlled entirely by the random seed.

Scenario	Mechanism	Expected outcome
Favourable seed	All centroids start near distinct true cluster centres	Algorithm converges quickly to a near-optimal solution → high silhouette
Unfavourable seed	Two or more centroids start inside the same true cluster	The algorithm converges to a local minimum where one true cluster is split and another is merged → lower silhouette

Practical implication: The spread in silhouette scores across seeds (with `n_init=1`) reveals how sensitive the dataset is to initialisation. A small spread means the clusters are well-separated - most seeds lead to the same solution. A large spread means the cluster structure is ambiguous or the classes overlap significantly.

The production default of `n_init=10` (or higher) addresses this by selecting the best result from multiple runs, making the final clustering robust to any single bad seed.



Summary

Parameter varied	Type of effect	Impact on silhouette
k (number of clusters)	Structural	Too few → groups merged, score drops. Too many → groups fragmented, score drops. Peak ≈ natural cluster count.
Random seed (initialisation)	Stochastic	With <code>n_init=1</code> , bad seeds trap the model in local optima. Variance across

Parameter varied	Type of effect	Impact on silhouette
		seeds measures how ambiguous the cluster structure is.

Together, these two investigations illustrate the two main sources of instability in KMeans: the **model complexity** (k) and the **optimisation landscape** (initialisation). In practice, both are addressed systematically - k by validation metrics, initialisation by `n_init > 1`.

Generated on 2026-05-14 | Mosudi I. O., FH Technikum Wien